# Lab 2 Design Problem

Encrypted Ramdisk

 $\begin{array}{c} \text{Keaton Boyle} -- 103\text{--}882\text{--}791 \\ \text{Anthony Ortega} -- 503\text{--}953\text{--}545 \end{array}$ 

February 21, 2013

# Introduction

This project provides functionality for users and applications to encrypt and decrypt a ramdisk using a key, and to read from and write to the resulting ramdisk when the correct key is provided. Certain functions are made available to userspace applications via the ecsprd.h header file. Since we're using the encryption algorithms provided by the kernel, we intend to support all of the block cipher encryptions that Linux already handles.

# Interface

The goal of our interface is to provide a few simple functions for the user to allow them to establish a connection with an encrypted ramdisk and then conduct any further communication with the encrypted ramdisk using basic read() and write() system calls. These functions are provided in the eosprd.h header file and are summarized here.

```
eosprd_open()
```

eosprd\_open() is a wrapper function that should be used in place of a normal open() system call whenever opening an encrypted ramdisk. It simply calls a system open() on the ramdisk and then provides the password via a newly-defined ioctl() call, EOSPRDIOCOPEN. The following is the basic structure of the eosprd\_open() call, without any of the error checking.

```
static int eosprd_open(const char *pathname, int flags, char *
    key)
{
    int fd;
    fd = open(pathname, flags);
    ... on failure, return -1 ...

    ioctl(fd, EOSPRDIOCOPEN, key);
    ... on failure, close fd and return -1 ...

    return fd;
}
```

If anything fails, -1 is returned. The call to open() can fail for any of the normal reasons, the ioctl() call will return -1 if the provided key does not match the key that the ramdisk was encrypted with (errno will be set to EKEYREJECTED). Additionally, if a ramdisk is not encrypted, the key argument should be NULL, or else the ioctl() call will return -1 and set errno to ENOSYS. (In such a use case, a simple open() could be used in place of eosprd\_open().) If the ioctl() does fail for any reason, the newly-opened file descriptor is closed before eoprd\_open() returns.

If all is successful, the newly opened file descriptor will be returned and the calling process will be able to read() and write() to it using normal system calls (according to the specifed flags). These calls will transparently encrypt and decrypt data coming to and from the ramdisk without the calling process ever needing to provide the key again, until the process exits or the file descriptor is closed.

### eopsrd\_encrypt()

eosprd\_encrypt() allows the user to encrypt, decrypt, or change the encryption key or scheme of a ramdisk. It takes four arguments: the path to the ramdisk file, the key the ramdisk was previously encrypted with (or NULL if the ramdisk was previously unencrypted), the new key to encrypt the ramdisk with (or NULL if the ramdisk should be unencrypted), and the name of the algorithm to be used to perform the encryption (one of the supported blocking encryption algorithms that Linux provides). The function is primarily a wrapper for a new iotcl() call, EOSPRDIOCENCRYPT, which is called with the function's argument packed into a struct encrypt\_args structure. The following is the basic structure of the eosprd\_encrypt() call, without some of the specifics.

```
static int eopsrd_encrypt(const char *pathname, char *oldkey,
   char *newkey, char *algo)
        int fd;
        struct encrypt_args e_args =
                 .oldkey = oldkey,
                 .newkey = newkey,
                 .algo = algo
        };
        fd = open(pathname, O_RDWR);
        ... on failure, return -1 ...
        ioctl(fd, EOSPRDIOCENCRYPT, &e_args);
        ... on failure, return -1 ...
        close(fd);
        ... on close failure, exit fatally- this leaves the
           ramdisk in a strange state ...
        return 0;
}
```

If the encryption or decryption is successful, 0 is returned. In case of failure, -1 is returned. Failure can occur for any of the typical reasons associated with open(), or for problems with the ioctl() call as discussed in eosprd\_open(), or if the algo argument isn't recognized (errno will be set to ENOSYS), or if another process has the the file open for encryption, reading, or writing (errno will be set to EBUSY).

This function, if successful, decrypts (if necessary) the ramdisk using the key oldkey and whatever scheme it was previously encrypted with and then re-encrypts it (if newkey is non-NULL) with the key newkey using the encryption algorithm specified by algo, or the previously-used algorithm if algo is NULL. If algo is NULL, newkey is non-NULL, and the ramdisk is coming from a non-encrypted state, the ioctl call will fail with errno set to ENOSYS.

In order to avoid race conditions and the possibility of one process blocking another's read or write by changing the encryption key, etc., this function will only succeed when no other process is reading, writing, or trying to encrypt/decrypt the ramdisk.

#### Notes:

If the user wishes to encrypt a file for the first time, oldkey should be NULL.

If the user wishes to decrypt a file, newkey should be NULL, and algo will be ignored. (This is also available via the eosprd\_decrypt() function.)

If the user wishes to change the key used to encrypt a file without changing the encryption algorithm, algo may be left NULL.

If the user wishes to switch encryption schemes without changing the keys, oldkey and newkey must be the same, and algo should be the new encryption scheme.

# eopsrd\_decrypt()

eosprd\_encrypt() is a simple inline function that simply calls eosprd\_encrypt() with the correct arguments for decryption. The entire function is shown here.

```
static inline int eosprd_decrypt(const char *pathname, char *
    oldkey)
{
    return eosprd_encrypt(pathname, oldkey, NULL, NULL);
}
```

#### Sample Code

The following is a simple sample application that encrypts a ramdisk, opens it, reads and writes to it, and closes it. Note that simple read(), write(), and close() calls are all that is necessary if the proper eosprd\_...() calls have been used ahead of time.

```
#include <unistd.h>
#include "eosprd.h"
int main(int nargs, char **argv)
{
        char buf [128];
        int fd;
        // encrypt the /dev/osprda ramdisk with key "password"
        // and the "blowfish" encryption algorithm
        eosprd_encrypt("/dev/osprda", NULL, "password", "
           blowfish");
        // open the file using the encryption key
        fd = eosprd_open("/dev/osprda", O_RDWR, "password");
        write(fd, "foo", 3);
        // foo will be encrypted and placed in /dev/osprda
        read(fd, buf, 3); // buf will now contain foo
        close(fd);
        // other processes will not be able to read "foo"
```

```
// without knowing the key
return 0;
}
```

# **Implementation**

The internal readings and writings to the ramdisk are done using largely the same code as in the bulk of Lab 2. The difference is that when encryption is enabled, we'll call kernel encryption and decryption functions before actually sending data to or from the ramdisk. Additionally, we'll need to keep track of if a ramdisk is encrypted, we'll need to know what encryption algorithm was used, we'll need some way to check if the user has provided the right key, and we'll need the file operations to somehow "remember" the key that the opening process provided, so that the user can use simple read() and write() calls without having to provide the key again. Furthermore, we need to be careful with concurrency issues, so that no process can reencrypt a file while another is trying to read it, or something of that sort. And of course, we'll actually need to do the encryption.

# Keeping track of encryption state

Solving this challenge is simple enough. We've added a flag, int encrypted, to the osprd\_info\_t structure. It is initially set to 0 when the ramdisk is first created. Later calls to eosprd\_encrypt() with non-NULL newkey values will set this flag to 1 so that all further calls to eosprd\_encrypt() and eosprd\_open() check for the correct key. Note that basic open() and read() calls (without first calling eosprd\_open()) will still succeed, but they will be reading the encrypted garbage data. A basic write() call (without first calling eosprd\_open()) will fail to write and return 0.

Calls to eosprd\_encrypt() with a NULL newkey (i.e., calls to eosprd\_decrypt()), if they provide the correct oldkey, will set this flag back to 0.

#### Keeping track of encryption algorithm

This too is fairly simple. Calls to eosprd\_encrypt() will use the newly-added char \*algo field of the osprd\_info\_t structure to retrieve and store the name of the algorithm used for encryption. This is initialized to NULL at the creation of the ramdisk.

#### Checking key correctness

We want some way to tell the user whether or not the correct key has been provided, but storing it somewhere in plaintext and checking against that seems to defeat some of the purpose of encryption. Instead, when the key is first provided by the eopsrd\_encrypt() function, we'll encrypt it with itself and store the result in the newly-added char \*key field of osprd\_info\_t. Then, when a user attempts to open the encrypted ramdisk, we'll decrypt key with the provided password, and if they match we'll be in the clear. Otherwise, the eosprd\_open() attempt will fail. This does leave access to the key user-side, but it avoids storing it in some fixed location.

# Remembering the key that a process has provided

Again, we'd rather not have a plaintext list of processes with priviledged access sitting around near our ramdisk, nor would we want to store the keys they've provided alongside them. But we need some way of remembering which open file descriptors have provided keys so that read()s and write()s on those file descriptors can use those keys. To do this, we'll store a pointer to the key in the f\_security field of the struct file attached to the open file descriptor. This way, every time read() or write() is called, we can refer to the f\_security field to do encryption and decryption. This information will disappear with the closing of the file descriptor.

#### Concurrency issues

Encryption or decryption should only occur when there are no other file descriptors open on a ramdisk. To guarantee this, encryption operations will always try to grab write locks on the ramdisk and will fail immediately if they are unable to. This isn't perfect, perhaps blocking would be better, but generally the intention is that initial encryption happens very rarely and should never be in parallel. Subsequent write() call encryption will occur transparently and reliably using already-implemented locking mechanisms.

# Encryption

Encryption will follow the basic encryption schemes provided by the kernel's crypto functions and will decrypt and encrypt one sector at a time.

#### Results

We'll have more information on our results once this has all been put to code, but generally we expect this to be a decently-working solution. It's unfortunate that keys, at some point or another, will be floating around in plaintext, but that's tough to avoid. The interactions and subtleties between user and kernel code is always a challenge, but we're working through it and trying to rely as much as possible on the pre-existing Lab 2 interfaces.

#### Work Division

Again, we'll have more information when we're finished, but design was worked out together by Mr. Boyle and Mr. Ortega. Most of the design document was written by Mr. Boyle, Mr. Ortega integrated the password interface into the osprdaccess program, and the rest of the coding is to be determined.